

## Implementation of IF in Cooma

# Table of Contents

- Model
  - Subtype Relation
  - Contra-variance and Co-variance
- Implementation
  - Syntax
  - Semantic Analysis
    - Subtype
    - Record and Variant Checks
    - Primitive Operations
  - Basic Types
  - Records and Variants
    - Design Notes
  - Capabilities
    - Design Notes
  - Other Additions

# Model

Information flow is modelled using the idea of information at two levels: public and secret. Secret variables (containing secret information) are denoted with a trailing `!`. For example, `var x : Int! = 1` declares a variable `x` of type `Int` which is classified as secret.

This notion extends to function arguments, return types, record fields and variant options. For example the following function `iden` takes a secret string as it's only argument and returns it:

```
1 def iden(x : String!) String! = x
```

We can't leak the value of the secret string using the following alternate version of `iden` as it is ill-typed:

```
1 def iden(x : String!) String = x
2                                     ^
3 error : Expected a public String but got a secret String
```

When extending this idea to records and variants, we are faced with a choice between allowing the whole type itself to also be classified as secret or keeping them as strictly public types (see [Records and Variants](#) for further discussion on design choices). In the current implementation, we have chosen to allow records and variants types to also be classified secret. With this decision comes additional caveats to the way records and variants may be declared:

- Public records and variants can contain a mixture of public and secret fields.
  - Secret records and variants can **ONLY** contain secret fields.
    - Functions within secret records must have both **ONLY** secret arguments and a secret return type.
- Enforcing argument types to be secret doesn't add any further restriction on what can be passed to them since public data can be classified as private. The secret return type does however enforce that whatever the function produces be secret. These rules extend to functions that may be returned from functions.

```

1 // Records
2 val mixedRec : { x : Int, y : Int! }
3 val secretRec : { x : Int!, y : Int! }!
4
5 // Variants
6 val mixedVar : < x : Int!, y : Int >
7 val secretVar : < x : Int!, y : Int! >!
8
9 // Secrets can't contain public fields
10 val errorRec : { x : Int }!
11      |-----|
12 var errorVar : < x : Int >!
13      |-----|
14 error : Secret record/variant can't contain public fields

```

Since the underlying representation of capabilities (`Reader/Writer/ReaderWriter`) are as records, the same rules apply then those of records. However, due to the fact that these capabilities are pre-defined, we have slightly more control over their underlying representation. We chose to implement them as secret records (see [Capabilities](#) for further discussion on design choices). An additional restriction is that either the data they consume (`Writer`) or produce (`Reader`) must also be secret to abide by the rules of secret records. This means the `ReaderWriter` does not have differing information levels for its `read` and `write` methods. This is intended to be the case as it makes more sense that a single capability exists at either a public or secret level, not a mix. If a mix is required, this can be done by declaring separate `Reader` and `Writer` capabilities. The underlying type of secret capabilities are:

```

1 type Reader! = { read : () => String! }!
2 type Writer! = { write : String! => () }!
3 type ReaderWriter! = { read : () => String!, write : String! => () }!

```

## Subtype Relation

We use the subtyping relation to enforce the notion of a "no write-down" policy. This does not allow secret variables to be used in place of public variables. Since we don't use any additional construct, other than `!`, the subtype is sufficient to enforce this. The relation  $T <: T!$  states that a type  $T$  is a subtype of  $T!$ , it's secret alternate. This allows public values to be made secret in order to perform computations with secret values, while disallowing secret values to become public.

## Contravariance and Covariance

The following describes the relationship between subtypes and arguments (contravariance) and subtypes and return types (covariance). It states that given a function  $S_1 \rightarrow S_2$  we can provide an argument  $T_1$  such that  $T_1$  is a subtype of the argument type  $S_1$ . The second part states that we can return an value of type  $T_2$  provided that it

is a supertype of the expected return type  $S_2$ .

$$\frac{T_1 <: S_1 \quad S_2 <: T_2}{S_1 \rightarrow S_2 <: T_1 \rightarrow T_2}$$

The below example was adapted from [Covariance and Contravariance \(computer science\)](#).

```
1 /*
2 Assumes the existence of the following relation Cat <: Animal
3 */
4
5 void contra(Animal a){ ... }
6 contra(new Cat()); // Ok - T1 <: S1 (Cat <: Animal)
7
8 Animal co(Cat a){
9     return a;      // Ok - S2 <: T2 (Cat <: Animal)
10 }
```

Using our subtype relation we can sketch some examples showing that we can classify public information but cannot declassify secrets:

```
1 // Using subtype relation Int <: Int! as an example
2
3 // Contravariance
4 {
5     def contra(x : Int!) ...
6     val x : Int = 10
7     contra(x)      // Ok - Int <: Int!
8 }
9 {
10    def contra(x : Int) ...
11    val x : Int! = 10
12    contra(x)      // Error - Int! <\: Int
13 }
14
15 // Covariance
16 {
17     def co(x : Int) Int! = x
18     val x : Int = 10
19     co(x)          // Ok - Int <: Int!
20 }
21 {
22     def co(x : Int!) Int = x
23     val x : Int! = 10
24     co(x)          // Error - Int <\: Int
25 }
```

## Implementation

This section will provide various implementation details of our information flow model in Cooma.

# Syntax

Currently in the Cooma parser, all types are hard coded as expressions. This affords us two ways of implementing secret variants of these types:

- Add an analogous SecT for every type.
- Group types together and add another higher level Sec type:

```
1 Expression {paren} =
2   ...
3   | PredefType '!'           {SecT}
4   | PredefType
5   ...
6
7 PredefType : Expression =
8   'Unit'                     {UnitT}
9   | '{' nest(FieldType ++ ',') \n '}' {RecT}
10  | '<' nest(FieldType ++ ',') \n '>' {VarT}
11  | 'Boolean'                 {BoolT}
12  | 'Int'                     {IntT}
13  | 'Reader'                  {ReaderT}
14  | 'ReaderWriter'            {ReaderWriterT}
15  | 'Writer'                  {WriterT}
16  | 'String'                  {StrT}
17  | 'Type'                    {TypT}
```

We have implemented the syntax for secret types using the second alternative in the `Cooma.syntax` file. This means that a the secret type is just a special wrapper around existing types. The form of secret types is this `SecT(Type)`. This is in contrast to an analagous type such as `SecIntT`. We have excluded function types `FunT` from the list as we do not feel that it's necessary for whole functions to be classified as secret. We feel that it's sufficient enough to only deal with the data being passed into the function.

## Semantic Analysis

This section is used to collect various implementataion notes on `SemanticAnalyser.scala`.

### Subtype

We can add our subtype realtion, which enforces our information flow, to the `subtype` function. Since secret types are wrappers over normal types, we can easily identify when we are checking some type `t` against a secret type `u` and pull the underlying type `s` out for comparison.

```
1 def subtype(t : Expression, u : Expression) : Boolean =
2   ...
3   case (_, SecT(s)) =>
4     subtype(t, s)
5   ...
```

### Record and Variant Checks

Due to our design choice to allow records or variants to be classified as secret, we need to add additional semantic

checks for `SecT(RecT)` and `SecT(VarT)`. The check needs to confirm that every field in the corresponding record or variant is of the form `SecT(␣)`. If a field is public (not wrapped with a `SecT`) it will be an error.

```
1 def checkSecretFields(fs : Vector[FieldType], t : Expression) : Messages = {
2   fs.flatMap(f => f.expression match {
3     case SecT(_) => noMessages
4     case _       => error(f, s"public field ${f.identifier} found in secret ${show(t)}")
5   })
6 }
```

An interesting case arises when a record contains a field which is also a record that is classified as public with **only** secret fields `val x : { a : { b : Int! } }!`. In this case, whilst all fields in `r` must be secret due to the trailing `!`, `a` is a public field. However, since all of the fields in `a` are secret, `a` is a direct subtype of `{ b : Int! }`. Currently, this is classified as an ill-typed secret record, but may be allowed with more research on the potential side-effects or this use. We could either wrap these records in a secret type (`SecT`) or perform an additional check in the `checkSecRecord` function:

```
1 def checkSecRecord(e : Expression) : Messages =
2   ...
3   case RecT(fs) =>
4     fs.forall {
5       case SecT(r : RecT) =>
6         checkSecRecord(r)
7       case SecT(_) => true
8       case _ => false
9     }
10  ...
```

Actually it turns out that nested fields are captured by the original check since it's performed on all `SecT(RecT)` and `SecT(VarT)` AST nodes. Since all fields must be public, when checking inner-records, it's not possible to have a public field of the form `SecT` and thus is safe to not add additional checks.

## Primitive Operations

We also need to consider the additional changes that we need to make to primitive operations. Our model dictates that we should be able to perform primitive operations between secret and public data as long as the result is also secret i.e. `Int! + Int = Int!`. This would mean special expectations need to be made for primitive operations when dealing with an operand of type `SecT`.

## Aliasing

Due to the way capabilities (`Reader/Writer/ReaderWriter`) are aliased, we need to add corresponding secret capabilities to the `alias` function, corresponding to their underlying implementation in `SymbolTable.scala`. The alias of the secret capabilities are the same as their public counterparts, prefixed with 'sec':

```
1 def alias(e : Expression) : Expression =
2   ...
3   case `secReaderT`      => SecT(ReaderT())
4   case `secReaderWriterT` => SecT(ReaderWriterT())
```

```

5 case `secWriterT`      => SecT(WriterT())
6 ...

```

This means we also need to add corresponding cases to `unalias` to allow us to 'move' the other way:

```

1 def unalias(n : ASTNode, t : Expression) : Option[Expression] =
2   ...
3   case SecT(ReaderT())      => Some(secReaderT)
4   case SecT(ReaderWriterT()) => Some(secReaderWriterT)
5   case SecT(WriterT())      => Some(secWriterT)
6   ...

```

## Basic Types

This section contains notes on the implementation of `Int`, `String` and `Boolean`.

### Int and String

No additional changes were needed other than the addition of the subtyping relation.

### Boolean

Due to the underlying implementation of `Boolean` the following additional were necessary:

- An additional case added to `subtype` function. This is due to the fact that `boolT` represents a boolean which unfolds to the variant `< False = (), True = () >`. Therefore when checking if a `boolT <: SecT(BoolT())` we need to catch this case and instead check `subtype(BoolT(), s)`. This is due to the fact that we cannot get the alias for `boolT` since the alias and `unalias` methods do not exist in the `SemanticAnalysis` object.

## Records and Variants

Only slight modifications, mainly to do with field selection on secret records were necessary. They include:

- Added extra case to the `Sel` case of the `type` function. This is necessary to make sure that when checking for a record in a `SecT(RecT())` that the actual `RecT()` AST node is checked instead of throwing an error since `SecT` doesn't contain the records fields.
- Added an extra case in the `checkFieldUse` function so that if we have a secret record `SecT(RecT())`, the fields of the `RecT()` are extracted and checked instead of throwing a "selection of X field from non-record type Y" error.

### Design Notes

The discussion here centers whether or not we should allow records or variants to be considered secret i.e. `val x : { a : Int! }!` and what that means in terms of additional checks for a program. As of its current implementation, Cooma allows records and variants to be defined as secret (as opposed only to fields). We feel this makes more sense than the alternative as records and variants themselves can be types. It also makes sense that a record or variant may be considered as secret as to limit the places that it can flow to. It also allows for a more idiomatic implementation of the built-in capabilities. We believe this decision makes the secret/public typing system slightly more expressive than its alternate option. The additional caveat of this however, is that additional checks must be in place, namely for secret records (`SecT(RecT())`) and secret variants (`SecT(VarT())`). We

restrict secret records to containing **only** secret fields as we feel this best fits with the ideas of a secret record. The same rule also applies to variants.

## Capabilities

Building on the changes required by records and variants, capabilities required a number of additions, namely:

- Separate secret capability variants defined in `SymbolTable.scala`. These are similar to their public counterparts but are prefixed with `sec`. The only changes from the public versions are:
  - The record is now wrapped in a `SecT` AST node.
  - The argument and return types for the pre-defined functions (`read` and `write`), excluding `Unit` have been wrapped in a `SecT` AST node also. This means that now secret capabilities only consume (`Writer.write`) and produce (`Reader.read`) secret values.
- Extra alias cases were added to the `alias` and `unalias` function as outlined in [Aliasing](#).

## Design Notes

Much of the previous discussion on public vs secret records and variant applies to capabilities also, since their underlying representation is as a record. This is covered in [Records and Variants - Design Notes](#). Another topic of interest is whether the `Unit` types in the `read` and `write` functions should be typed as `SecT(Unit)` or just `Unit`. Currently we are not convinced it makes much of a difference since we shouldn't need to worry since we don't need to check subtypes between capabilities i.e. only `aReader` can be used in place of a `Reader`, not another capability. Since these capabilities can only be assigned by the run-time (users cannot create them), we don't need to worry about user defined ones either. Currently said function only use the `Unit` but a shift to `SecT(Unit())` would be easy if needed.

## Other Additions

This section contains a number of additional changes.

An additional change was required in the `type` function to determine the type of `SecT` nodes:

```
1 lazy val type : Expression => Option[Expression]
2   ...
3   case _ : SecT =>
4     Some(TypT())
5   ...
```